

# CompStat 1

## Problem A

1)


Here using the inversion method, we get:

$$x = F^{-1}(u) = -\frac{1}{\lambda} \log(u), \quad u \sim U(0,1).$$

Thus we get xponential samples, by sampling from the uniform distribution, and use the inversion method. Under there is a function that generates a vector of  $n$  samples from an exponentially distributed random variable with rate  $\lambda$ .

```
# Generate samples from exponential distribution
rate <- 3 #Rate parameter
n <- 1000 #Number of samples
generate.exponential <- function(rate, n){
  u <- runif(n)
  x <- - 1/rate * log(u)
  return(x)
}
x <- generate.exponential(rate, n)
hist(x, density=20, breaks = 30, prob=TRUE)
curve(dexp(x, rate=rate), add=TRUE, col="darkblue", lwd=2)
```

**Histogram of x**



2)

The cdf  $G(x)$  of  $g(x)$  and its inverse  $G^{-1}$  are


$$G(x) = \begin{cases} \int_0^1 cx^{\alpha-1}dx & , 0 < x \leq 1 \\ \int_1^\infty ce^x dx & , 1 \leq x \\ 0 & , x \leq 0 \end{cases} = \begin{cases} c(\alpha-1)x^\alpha & , 0 < x \leq 1 \\ c(e^{-1} - e^{-x}) & , 1 \leq x \\ 0 & , x \leq 0 \end{cases}$$

$$G^{-1}(u) = \begin{cases} \left(\frac{u}{c(\alpha-1)}\right)^{\frac{1}{\alpha}} & , \text{SJEKK GRENSER } 0 < \left(\frac{u}{c(\alpha-1)}\right)^{\frac{1}{\alpha}} \leq 1 \\ -\log(e^{-1} - \frac{u}{c}) & , 1 \leq -\log(e^{-1} - \frac{u}{c}) \\ 0 & , u \leq 0 \end{cases}$$

Writing a function that returns  $g$ -samples.

```
# Generate samples from g(X)
rate <- 3 #Rate parameter
n <- 1000 #Number of samples
alpha <- 0.1
generate.samples.g<- function(alpha, n){
  c <- alpha*exp(1)/(alpha+exp(1))
  u <- runif(n, 0,1)
  x <- c()
  for (i in 1:n){
    if (u[i] < c/(alpha)){
      x <- c(x, (u[i]*(alpha)/c)^(1/alpha))
    }
    else{
      x <- c(x, 1-log(exp(1)/c*(1-u[i])))
    }
  }
  return(x)
}
x <- generate.samples.g(alpha, n)
hist(x, density=20, breaks = 20, prob=TRUE)
```

## Histogram of x




The histogram is used to check that `generate.samples.g` returns the correct value. As we see here, ...


### 3)

The function `generate.normal` uses Box-Muller to generate an  $n$ -vector of independent samples from the standard normal distribution.

```
# Generate samples from normal distribution with Box-Muller
n <- 10000 #Number of samples
generate.normal <- function(n){
  u1 <- runif(n)
  u2 <- runif(n)
  x1 <- sqrt(-2*log(u1))*cos(2*pi*u2)
  x2 <- sqrt(-2*log(u1))*sin(2*pi*u2)
  X <- rbind(x1, x2)
  return (X)
}
X <- generate.normal(n)
hist(X[1,], density=20, breaks = 30, prob=TRUE)
curve(dnorm, add=TRUE, col="darkblue", lwd=2)
```



```
hist(X[2,], density=20, breaks = 30, prob=TRUE)  
curve(dnorm, add=TRUE, col="darkblue", lwd=2)
```



```
## [1] "Observed: "
## [1] -0.004325752

##           x1           x2
## x1 0.9935009701 0.0002658448
## x2 0.0002658448 1.0021586764
```

As we see, the observed mean and variance is fairly close to the moments of the standard normal distribution, which are zero and one. The empirical covariance  $Cov(x_1, x_2)$  is also close to zero, as expected.

4)

Writing function `generate.d.normal` that generates samples from a  $d$ -variate normal distribution with mean vector  $\mu$  and covariance matrix  $\Sigma$ .

```
# Generate n samples from d variate normal distribution
n <- 10000 #Number of samples
mu <- c(0,0)
sigma <- cbind(c(1,0.5),c(0.5,1))
generate.d.normal <- function(n, mu, sigma){
  A <- chol(sigma)
  d <- length(mu)
  X <- generate.normal(n*d)
  x <- matrix(X[1,], nrow=d, ncol=n)
  Z <- mu+t(A)%*%x
}
Z <- generate.d.normal(n, mu, sigma)
```

```
## [1] "Expected: "
## [1] 0 0
##      [,1] [,2]
## [1,]  1.0  0.5
## [2,]  0.5  1.0
## [1] "Observed: "
## [1] -0.0005811538
##      [,1]      [,2]
## [1,] 1.0195376 0.5104683
## [2,] 0.5104683 1.0053814
```

As we see,  $\mu$  is close to ER MU KUN 2-VARIAT?

## B

```
#Function f(x)
function.f <- function(x, alpha){
  if (x>0){
    return (x^(alpha-1)*exp(-x)/gamma(alpha))
  }
  else{
    return(0)
  }
}

#Function g(x)
function.g <- function(x, alpha){
  c = alpha*exp(1)/(alpha+exp(1))
  if (x>0 && x<1){
    return (c*x^(alpha-1))
  }
  else if(x >= 1){
    return(c*exp(-x))
  }
  else{
    return(0)
  }
}

# Generate n samples from d variate normal distribution
n <- 1000 #Number of samples
alpha <- 0.3
rejection.sampling <- function(n, alpha){
  c <- (alpha+exp(1))/(gamma(alpha)*alpha*exp(1))
  samples <- c()
  for (i in 1:n){
    repeat{
      x <- generate.samples.g(alpha, 1)
      acceptance.prob <- function.f(x, alpha)/(c*function.g(x, alpha))
      if(runif(1) <= acceptance.prob){
        samples <- c(samples, x)
      }
    }
  }
}
```

```

        break
    }
}
return (samples)
}
y <- rejection.sampling(n, alpha)

## [1] "Expected: "
## Mean: 0.3
## Variance: 0.3
## [1] "Observed: "
## Mean: 0.2925686
## Variance: 0.2867875

# An analytic log(gamma(x, alpha))-function
log.gamma <- function(x, alpha){
  return((alpha-1)*log(x)-x)
}

# Ratio-of-uniforms method
n <- 1000
alpha <- 30
ratio.of.uniforms <- function(n, alpha){
  sample.counter <- 0
  b.minus <- 0
  log.b.plus <- (alpha+1)/2*(log(alpha+1)-1)
  log.a <- (alpha-1)/2*(log(alpha-1)-1)
  samples <- c()
  for(i in 1:n){
    repeat{
      sample.counter <- sample.counter + 1
      log.x1 <- -generate.exponential(1, rate=1) + log.a
      log.x2 <- -generate.exponential(1, rate=1) + log.b.plus
      log.y <- log.x2 - log.x1
      y <- exp(log.y)
      if (2*log.x1 <= log.gamma(y, alpha)){
        samples <- c(samples, y)
        break
      }
    }
  }
  cat("Rejection ratio: ", (sample.counter-n)/n , "\n")
  return (samples)
}
y <- ratio.of.uniforms(n, alpha)

## Rejection ratio: 3.458
## [1] "Expected: "
## Mean: 30
## Variance: 30

```

```
## [1] "Observed: "
## Mean: 30.00993
## Variance: 30.18444
n = 1000
alpha = 0.1
beta = 2
sample.gamma <- function(n, alpha, beta){
  if(alpha < 1){
    return(rejection.sampling(n, alpha)/beta)
  }
  else if (alpha == 1){
    return(generate.exponential(1, n)/beta)
  }
  else if (alpha > 1){
    return(ratio.of.uniforms(n, alpha)/beta)
  }
  else{
    return(0)
  }
}
y <- sample.gamma(n, alpha, beta)

## [1] "Expected: "
## Mean: 0.05
## Variance: 0.025
## [1] "Observed: "
## Mean: 0.04668849
## Variance: 0.01995032
```

## Problem C

1)

Claim: Given independent variables  $z_k \sim \text{gamma}(\alpha_k, \beta = 1)$  for  $k = 1, \dots, K$ , define  $x_k = z_k \cdot (\sum_{k=1}^K z_k)^{-1}$  for  $k = 1, \dots, K$ . Then vector  $(x_1, \dots, x_K)$  has the Dirichlet distribution with parameter vector  $\alpha = (\alpha_1, \dots, \alpha_K)$ .  
 Demonstration: Consider vector  $x = (x_1, \dots, x_{K-1})$  with  $x_k$  as before and variable  $v = \sum_{k=1}^K z_k$ . Express the transformation

$$Z_k = X_k \cdot V \text{ for } k = 1, \dots, K-1 ; Z_K = V \cdot X_K = V(1 - X_1 - \dots - X_{K-1})$$

$X, V$  now has joint distribution  $f_{X,V} = |J| \cdot f_Z(z) = |J| \cdot \prod_{k=1}^K e^{-z_k} z_k^{\alpha_k-1} \cdot \frac{1}{\Gamma(\alpha_k)}$

$$f_{X,V} = |J| \cdot f_Z(z) = |J| \cdot \prod_{k=1}^K e^{-z_k} z_k^{\alpha_k-1} \cdot \frac{1}{\Gamma(\alpha_k)}$$

with



$$J = \begin{bmatrix} V & 0 & \dots & X_1 \\ 0 & V & \dots & \vdots \\ \vdots & \vdots & \ddots & X_{K-1} \\ -V & \dots & -V & X_K \end{bmatrix} \sim \begin{bmatrix} V & 0 & \dots & X_1 \\ 0 & V & \dots & \vdots \\ \vdots & \vdots & \ddots & X_{K-1} \\ 0 & 0 & 0 & 1 \end{bmatrix} \implies |J| = V^{K-1}$$

Here we have used row reduction on the bottom row, along with the fact that  $X_K = \sum_{k=1}^{K-1} X_k$  and determinant rule for triangular matrices.

Isolate the  $v$ -part of  $f_{X,V}$  and integrate with respect to  $v$ , yielding the marginal distribution  $f_X$ .

$$f_{X,V} = v^{K-1} \prod_{k=1}^K e^{-x_k v} (x_k v)^{\alpha_k - 1} \cdot \frac{1}{\Gamma(\alpha_k)} = v^{\sum \alpha_k - 1} e^{-v} \prod_{k=1}^K x_k^{\alpha_k - 1} \cdot \frac{1}{\Gamma(\alpha_k)}$$

$$f_X(x) = \prod_{k=1}^K x_k^{\alpha_k - 1} \cdot \frac{1}{\Gamma(\alpha_k)} \cdot \int_0^\infty v^{\sum \alpha_k - 1} e^{-v} dv$$

$$f_X(x) = \left( \prod_{k=1}^K x_k^{\alpha_k - 1} \cdot \frac{1}{\Gamma(\alpha_k)} \right) \cdot \Gamma\left(\sum_{k=1}^K \alpha_k\right) = \frac{\Gamma\left(\sum_{k=1}^K \alpha_k\right)}{\prod_{k=1}^K \Gamma(\alpha_k)} \prod_{k=1}^{K-1} x_k^{\alpha_k - 1} \cdot \left(1 - \sum_{k=1}^{K-1} x_k\right)^{\alpha_K - 1}$$

where the integral resolves from the fact that the Gamma distribution with shape  $\sum_{k=1}^K \alpha_k$  and rate one integrates to one. We again used that  $X_K = 1 - \sum_{k=1}^{K-1} X_k$ . ## 2)

```
# Here, the user can manually change the alpha vector
# in order to draw a sample from any chosen Dirichlet distribution.
alphas.from.user=c(1.34534,0.9345,0.6356,20,14,5,2)

gammas.in.dirichlet <- function(alpha){
  n=length(alpha)
  gammas <- c()
  for(i in 1:n){
    gammas <- c(gammas, sample.gamma(1,alpha[i],1))
  }
  return(gammas)
}

generate.dirichlet <- function(alpha){ # Creates Dirichlet distributed vector
  g<-gammas.in.dirichlet(alpha)        # with the K chosen alphas
  print(g)                             # as K-variate parameter vector.
  v<-sum(g)
  z=g/v
  return(z)
}

dirich.var<-generate.dirichlet(alphas.from.user)

## Rejection ratio: 0
## Rejection ratio: 0
## Rejection ratio: 11
## Rejection ratio: 0
## Rejection ratio: 1
## [1] 0.06521793 0.51730071 0.41778437 13.86248879 13.92143278 9.03238379
## [7] 2.76539427
```

```

dirich.var           # The elements in the Dirichlet distribution

## [1] 0.001607065 0.012747047 0.010294819 0.341592033 0.343044499 0.222571169
## [7] 0.068143366

print(sum(dirich.var)) # The sum of the elements in the Dirichlet distribution

## [1] 1

```

In the code, the normalising of the Dirichlet distributed vector  $x$  occurs after the gammas  $z$  have been sampled, so  $\|x\|_1$  is always one, which is equal to the joint expected value of  $x$ 's elements.

## Problem D

1)

Given a prior Beta(1, 1), which is equal to the uniform distribution on (0, 1), and a multinomial density function:

$$f(\mathbf{y}|\theta) \propto (2 + \theta)^{y_1} (1 - \theta)^{y_2 + y_3} \theta^{y_4},$$

we have that the posterior is distributed with:

$$f(\theta, \mathbf{y}) \propto (2 + \theta)^{y_1} (1 - \theta)^{y_2 + y_3} \theta^{y_4}, \quad \theta \in (0, 1).$$


We also have the values for  $\mathbf{y}$ :  $y_1 = 125$ ,  $y_2 = 18$ ,  $y_3 = 20$ ,  $y_4 = 34$ . Thus we can use a rejection sample method, where we sample thetas from the prior, compute the acceptance probability, and then reject/include a random number into sample. To find the constant  $c$ , we have optimized over the posterior distribution. This will not be the real  $c$  value, but  $c$  times the proportional constant of the posterior distribution and the multinomial distribution. There is however no need to find this constant as when dividing the posterior by  $c$ , it would cancel.

```

# Rejection sampling
n <- 20000 #Number of samples
posterior.f <- function(theta){
  y <- c(125, 18, 20, 34)
  return ((2+theta)^y[1]*(1-theta)^(y[2]+y[3]))*theta^y[4])
}
rejection.sampling.multinomial <- function(n){
  sample.counter <- 0
  c <- optimize(posterior.f, c(0, 1), maximum=TRUE)$objective
  samples <- c()
  for(i in 1:n){
    repeat{
      sample.counter <- sample.counter + 1
      theta <- runif(1)
      acceptance.prob <- posterior.f(theta)/c
      u <- runif(1)
      if (u <= acceptance.prob){
        samples <- c(samples, theta)
        break
      }
    }
  }
  li <- list("samples" = samples, "number.samples" = sample.counter)
  return (li)
}

```

```
}
theta <- rejection.sampling.multinomial(n)
```



```
## Theta sample mean 0.6226123
```

Here is a plot over the density of the sampled  $\theta$ 's as well as the sampled mean. As we can see the sample mean of  $\theta$  is circa 0.623, which may suggest that the prior was not right (as we would expect mean close to 0.5 from a uniform (0, 1) distribution).

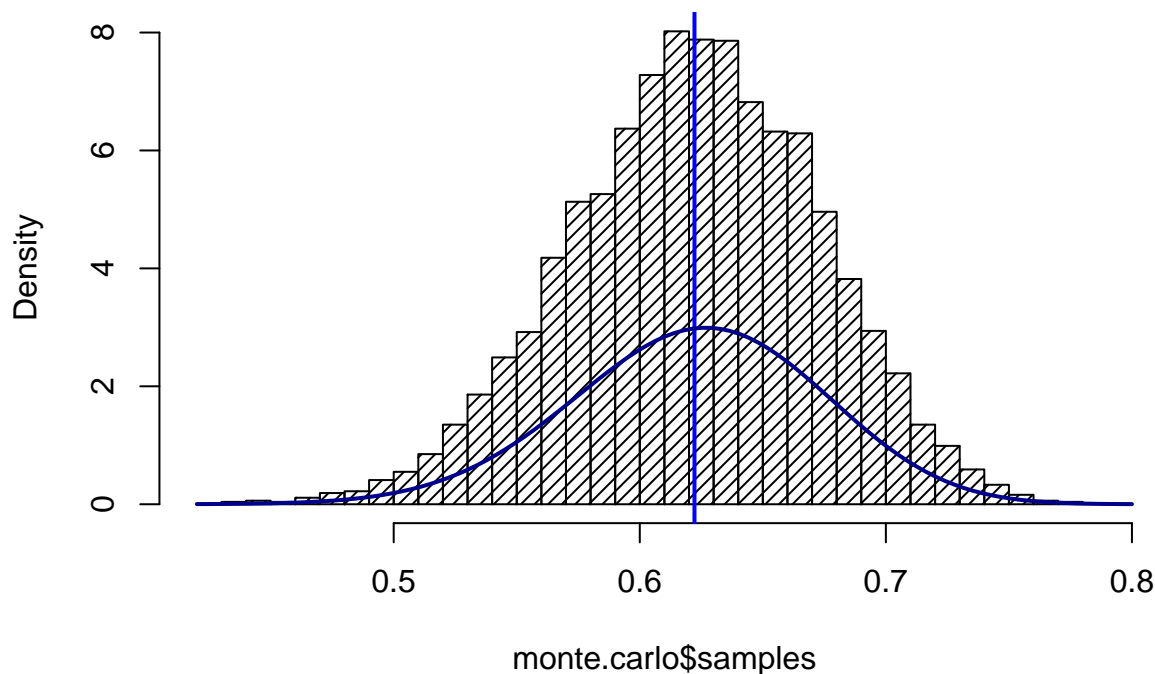
2)

Here we have

```
monte.carlo.integration <- function(){
  m <- 10000 # Number of samples
  samples <- rejection.sampling.multinomial(m)$samples
  mu <- sum(samples)/m
  li <- list("observed_mean" = mu, "samples" = samples)
  return (li)
}
monte.carlo <- monte.carlo.integration()

hist(monte.carlo$samples, density=20, breaks = 30, prob=TRUE)
abline(v = mean(monte.carlo$samples), col = "blue", lwd = 2)
curve(posterior.f(x)/mean(posterior.f(x)), add=TRUE, col="darkblue", lwd=2)
```

## Histogram of monte.carlo\$samples



```
numerical.integration <- function(theta){
  return (theta*posterior.f(theta))
}
Analytic_mean <- integrate(numerical.integration, lower=0,upper=1)$value/integrate(posterior.f, lower=0
```

```
## Observed mean: 0.6222184
```

```
## Analytic mean: 0.6228061
```

3)

Sampling using the rejection algorithm leads to a number of rejected samples. In this case we have counted the ratio as we sampled, marked as observed, as well as found the analytic solution which will be:

$$c^{-1} = \int_0^1 \frac{f(x)}{c} dx.$$

Using the built in optimize and integrate function, we get the integral, which will be the ratio  $\frac{\text{Total samples}}{\text{Accepted samples}}$ :

```
## Observed: 7.88975
```

```
## Analytic: 7.799308
```

As we see we have a observed ratio close to the analytic ratio.

4)

We have the samples from earlier which are sampled with the prior  $\text{Beta}(1, 1)$ :  $x_1, \dots, x_n \sim g(x)$ . However we can change this by using importance sampling:

$$w_i = \frac{f(x_i)}{g(x_i)}$$

Here  $f(x) \sim \text{Beta}(1, 5)$ . Since the density function will stay the same, all the elements will cancel, and we will get weights:

$$w_i = \frac{\text{Beta}(1, 5)}{\text{Beta}(1, 1)} = \text{Beta}(1, 5) = (1 - \theta)^4.$$

Thus the importance sampling algorithm will be for this case:

$$\tilde{\mu}_{IS} = \frac{\sum h(\theta_i)w(\theta_i)}{\sum w(\theta_i)} = \frac{\sum \theta_i w(\theta_i)}{\sum w(\theta_i)}, \quad w(\theta_i) = (1 - \theta_i)^4$$

Running this algorithm as code gives a estimator  $\tilde{\mu}_{IS}$  for the new mean.

```
beta.distribution <- function(x, alpha, beta){  
  return(x^(alpha-1)*(1-x)^(beta-1)/beta(alpha, beta))  
}  
importance.sampling <- function(n){  
  samples <- rejection.sampling.multinomial(n)$samples  
  w = (1-samples)^4  
  mu.is <- sum(samples*w)/sum(w)  
  return(mu.is)  
}  
  
importance.sampling(20000)
```

```
## [1] 0.5960066
```

Here we see from the print out, that the importance sampling shifts the observed mean from 0.623 to 0.596.